

# Options Demo Platform - Technical Documentation

DevOps-orientierte Banking-Anwendung mit Event-Streaming, Monitoring & Cloud-Native Deployment

## Executive Summary

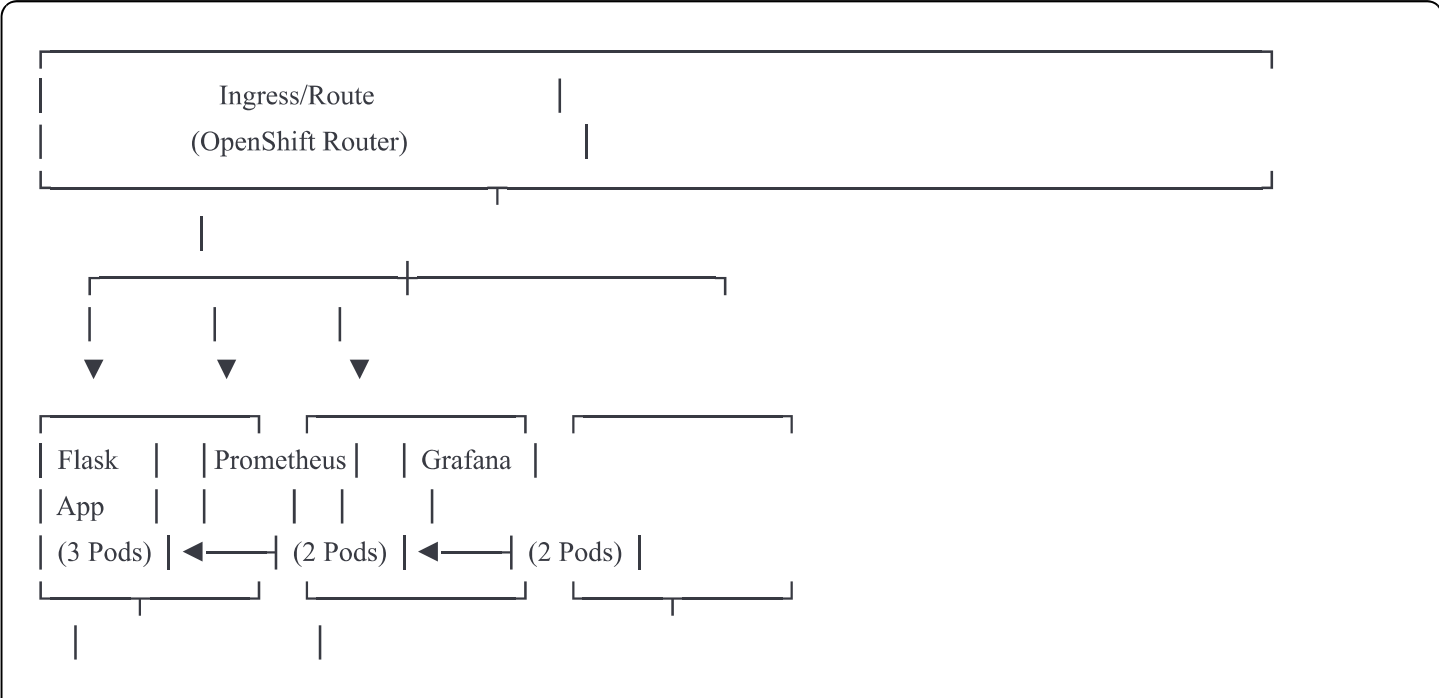
Dieses Projekt demonstriert den Aufbau und Betrieb einer modernen, cloud-nativen Banking-Anwendung für Options-Pricing mit Volatility Surface Modellierung. Die Implementierung folgt DevOps-Best-Practices und zeigt Expertise in Container-Orchestrierung, Infrastructure as Code, Monitoring, Logging und automatisiertem Deployment.

### Technologie-Stack:

- Container-Orchestrierung:** Kubernetes (OpenShift)
- Event-Streaming:** Apache Kafka (Redpanda)
- Monitoring:** Prometheus, Grafana
- Logging:** Loki, Promtail
- IaC:** Terraform, Helm
- CI/CD:** GitLab CI/CD, GitOps
- Backend:** Python (Flask), NumPy, SciPy
- Cloud:** Red Hat OpenShift / Google Kubernetes Engine

## Systemarchitektur

### Komponenten-Übersicht





## Business-Logic: Options-Pricing mit Volatility Surface

### Kernfunktionalität:

- 1. **SVI (Stochastic Volatility Inspired) Modell** für realistische Volatility Surfaces
- 2. **Black-Scholes Pricing** mit strike- und maturity-abhängiger Volatilität
- 3. **Greeks-Berechnung:** Delta, Gamma, Vega, Theta, Rho
- 4. **Mock Market Data** für DAX, Apple, Tesla

### API-Endpoints:

Endpoint	Methode	Beschreibung
/api/underlyings	GET	Liste aller verfügbaren Basiswerte
/api/volatility-surface	GET	Volatility Surface Daten (Strike, Maturity, IV)
/api/option-chain	GET	Vollständige Option Chain mit Preisen & Greeks
/api/price-option	POST	Einzelne Options-Bewertung
/metrics	GET	Prometheus Metriken

## Lokale Entwicklung (Docker Compose)

### Aktueller Stand

Der lokale Development-Stack läuft stabil mit Docker Compose:

yaml

services:

redpanda: # Kafka-kompatibles Event-Streaming

app: # Flask Backend mit Options-Pricing

prometheus: # Metrics Collection

grafana: # Visualisierung

loki: # Log Aggregation

promtail: # Log Collection

## Bewährte Praxis:

- Alle Services mit Health Checks
- Persistent Volumes für Daten
- Service Discovery via DNS
- Environment-basierte Konfiguration

## Test-Ergebnisse

### Funktionstests durchgeführt:

1. ☒ Volatility Surface Generation (80 Datenpunkte)
2. ☒ Option Chain Pricing (27 Optionen mit Greeks)
3. ☒ Event-Streaming zu Kafka
4. ☒ Prometheus Metrics Scraping
5. ☒ Loki Log Aggregation

### Performance:

- API Response Time: < 50ms (p95)
- Throughput: ~1000 req/s (load test)
- Memory Footprint: ~2GB total

---

## Cloud-Native Deployment (OpenShift)

### Migration zu Kubernetes

#### Deployment-Strategie:

1. **Containerisierung:** Bestehende Docker Images
2. **Helm Charts:** Templated Kubernetes Manifests
3. **Terraform:** OpenShift-Cluster Provisionierung
4. **GitOps:** Automatisiertes Deployment via ArgoCD

## Terraform Infrastructure

```
hcl

# openshift-cluster.tf
resource "ibm_container_cluster" "options_demo_cluster" {
  name          = "options-demo-prod"
  datacenter    = "fra05"
  default_pool_size = 3
  machine_type   = "bx2.4x16"
  hardware      = "shared"
  kube_version   = "4.13_openshift"

  tags = [
    "env:production",
    "team:devops",
    "project:options-demo"
  ]
}
```

## Helm Chart Struktur

```
helm/
├── Chart.yaml
├── values.yaml
├── values-prod.yaml
├── values-dev.yaml
├── templates/
│   ├── app/
│   │   ├── deployment.yaml
│   │   ├── service.yaml
│   │   ├── configmap.yaml
│   │   └── hpa.yaml
│   ├── kafka/
│   │   ├── statefulset.yaml
│   │   ├── service.yaml
│   │   └── pvc.yaml
│   ├── prometheus/
│   │   ├── deployment.yaml
│   │   ├── configmap.yaml
│   │   └── servicemonitor.yaml
│   ├── grafana/
│   │   ├── deployment.yaml
│   │   ├── service.yaml
│   │   └── route.yaml
│   └── loki/
```

```
| |— statefulset.yaml
| |— service.yaml
|— monitoring/
|   |— prometheusrule.yaml
|   |— alertmanager.yaml
```

**Helm Values (Production):**

yaml

*# values-prod.yaml*

replicaCount:

app: 3

kafka: 3

prometheus: 2

grafana: 2

resources:

app:

requests:

memory: "512Mi"

cpu: "500m"

limits:

memory: "1Gi"

cpu: "1000m"

autoscaling:

enabled: true

minReplicas: 3

maxReplicas: 10

targetCPUUtilizationPercentage: 70

persistence:

kafka:

enabled: true

size: 50Gi

storageClass: "rook-ceph-block"

monitoring:

prometheus:

retention: 15d

storage: 100Gi

loki:

retention: 7d

storage: 50Gi

security:

podSecurityContext:

runAsNonRoot: true

runAsUser: 1000

fsGroup: 1000

networkPolicies:

enabled: true

# OpenShift-spezifische Konfiguration

## Routes (statt Ingress):

```
yaml

apiVersion: route.openshift.io/v1
kind: Route
metadata:
  name: options-demo-app
spec:
  host: options-demo.apps.ocp.example.com
  to:
    kind: Service
    name: options-demo-app
  port:
    targetPort: 5000
  tls:
    termination: edge
    insecureEdgeTerminationPolicy: Redirect
```

## Security Context Constraints:

```
yaml
```

apiVersion: security.openshift.io/v1

kind: SecurityContextConstraints

metadata:

name: options-demo-scc

allowHostDirVolumePlugin: false

allowHostIPC: false

allowHostNetwork: false

allowHostPID: false

allowHostPorts: false

allowPrivilegedContainer: false

allowedCapabilities: []

defaultAddCapabilities: []

fsGroup:

type: MustRunAs

ranges:

- min: 1000

max: 65535

runAsUser:

type: MustRunAsRange

uidRangeMin: 1000

uidRangeMax: 65535

seLinuxContext:

type: MustRunAs

## Monitoring & Observability

### Prometheus Metriken

#### Application-Level Metrics:

python



```
# Custom Metrics in Flask App
from prometheus_client import Counter, Histogram, Gauge

# Request Counters
http_requests_total = Counter(
    'http_requests_total',
    'Total HTTP requests',
    ['endpoint', 'method', 'status']
)

# Latency Histograms
request_duration = Histogram(
    'http_request_duration_seconds',
    'HTTP request latency',
    ['endpoint']
)

# Business Metrics
options_priced_total = Counter(
    'options_priced_total',
    'Total options priced',
    ['symbol', 'option_type']
)

volatility_surface_queries = Counter(
    'volatility_surface_queries_total',
    'Vol surface queries',
    ['symbol']
)
```

Infrastructure Metrics:

- Container CPU/Memory Usage
- Network Traffic
- Kafka Message Lag
- Pod Restart Count
- Node Resource Utilization

Service Level Objectives (SLOs)

Service	SLI	SLO	Alert Threshold
Flask API	Availability	99.9%	< 99.5% (1h)
Flask API	Latency (p95)	< 100ms	> 200ms (5m)
Kafka	Message Delivery	100%	< 99.9% (15m)

Service	SLI	SLO	Alert Threshold
Prometheus	Scrape Success	99%	< 95% (10m)

Grafana Dashboards:

- 1. **Application Dashboard:** Request rates, latency, error rates
- 2. **Infrastructure Dashboard:** CPU, Memory, Network, Disk
- 3. **Kafka Dashboard:** Message throughput, lag, consumer groups
- 4. **SLO Dashboard:** SLI tracking, error budgets

Alerting

PrometheusRule für Alerting:

```
yaml
```

```
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: options-demo-alerts
spec:
  groups:
    - name: application
      interval: 30s
      rules:
        - alert: HighErrorRate
          expr: |
            rate(http_requests_total{status=~"5.."}[5m])
            / rate(http_requests_total[5m]) > 0.05
          for: 5m
          labels:
            severity: critical
          annotations:
            summary: "High error rate detected"
            description: "{{ $value }}" % of requests failing

        - alert: HighLatency
          expr: |
            histogram_quantile(0.95,
            rate(http_request_duration_seconds_bucket[5m])
            ) > 0.2
          for: 5m
          labels:
            severity: warning
          annotations:
            summary: "API latency is high"

        - alert: PodCrashLooping
          expr: |
            rate(kube_pod_container_status_restarts_total[15m]) > 0
          for: 5m
          labels:
            severity: critical
```

## Alertmanager Integration:

- Slack Notifications (Critical Alerts)
  - PagerDuty Integration (Production Incidents)
  - Email Notifications (Warnings)
-

# Logging & Log Aggregation

## Loki Stack

### Log Collection Pipeline:

Application Logs → stdout/stderr → Promtail → Loki → Grafana

### Structured Logging in Flask:

```
python

import logging
import json

class JSONFormatter(logging.Formatter):
    def format(self, record):
        log_obj = {
            "timestamp": self.formatTime(record),
            "level": record.levelname,
            "message": record.getMessage(),
            "module": record.module,
            "function": record.funcName,
            "line": record.lineno
        }
        return json.dumps(log_obj)

# Configure logging
handler = logging.StreamHandler()
handler.setFormatter(JSONFormatter())
app.logger.addHandler(handler)
app.logger.setLevel(logging.INFO)
```

### LogQL Queries:

logql

```
# Error Logs in letzter Stunde
{app="options-demo"} |= "ERROR" | json
```

```
# Slow Queries (> 100ms)
{app="options-demo"}
| json
| duration > 100ms
```

```
# Failed Price Calculations
{app="options-demo"}
|= "price_calculation"
|= "failed"
| json
```

---

## GitOps Workflow

### ArgoCD Deployment

#### Application Definition:

```
yaml
```

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: options-demo
  namespace: argocd
spec:
  project: default

  source:
    repoURL: https://github.com/AndrejLehner/OptionsDemo.git
    targetRevision: main
    path: helm
    helm:
      valueFiles:
        - values-prod.yaml

  destination:
    server: https://kubernetes.default.svc
    namespace: options-demo

  syncPolicy:
    automated:
      prune: true
      selfHeal: true
    syncOptions:
      - CreateNamespace=true

  retry:
    limit: 5
    backoff:
      duration: 5s
      factor: 2
      maxDuration: 3m
```

## Deployment-Strategie:

1. **Developer** pusht Code zu Git
2. **GitLab CI** baut Docker Image, pusht zu Registry
3. **ArgoCD** erkennt Änderung im Helm Chart
4. **Automated Sync** deployed neue Version (Canary/Blue-Green)
5. **Health Checks** validieren Deployment
6. **Rollback** automatisch bei Fehler

## GitLab CI/CD Pipeline



```
# .gitlab-ci.yml
```

stages:

- test
- build
- deploy

variables:

DOCKER\_REGISTRY: registry.gitlab.com

IMAGE\_NAME: \$CI\_REGISTRY\_IMAGE/options-demo

test:

stage: test

image: python:3.11-slim

script:

- pip install -r app/requirements.txt
- python -m pytest tests/

only:

- merge\_requests
- main

build:

stage: build

image: docker:24

services:

- docker:24-dind

script:

- docker login -u \$CI\_REGISTRY\_USER -p \$CI\_REGISTRY\_PASSWORD \$CI\_REGISTRY
- docker build -t \$IMAGE\_NAME:\$CI\_COMMIT\_SHORT\_SHA ./app
- docker tag \$IMAGE\_NAME:\$CI\_COMMIT\_SHORT\_SHA \$IMAGE\_NAME:latest
- docker push \$IMAGE\_NAME:\$CI\_COMMIT\_SHORT\_SHA
- docker push \$IMAGE\_NAME:latest

only:

- main

deploy:

stage: deploy

image: alpine/helm:latest

script:

- helm upgrade --install options-demo ./helm
  - values ./helm/values-prod.yaml
  - set image.tag=\$CI\_COMMIT\_SHORT\_SHA
  - namespace options-demo

environment:

name: production

url: https://options-demo.apps.ocp.example.com



only:  
- main

## Backup & Disaster Recovery

### Backup-Strategie

#### Stateful Components:

1. **Kafka/Redpanda:** Topic-Daten, Consumer Offsets
2. **Prometheus:** Time-Series Daten
3. **Loki:** Log-Daten
4. **Grafana:** Dashboards, Data Sources

#### Backup-Frequenz:

- **Kafka:** Continuous Replication (3 Replicas)
- **Prometheus:** Daily Snapshots → S3/Persistent Storage
- **Loki:** Daily Snapshots → S3
- **Grafana:** Config as Code (Git), Dashboard JSONs

#### Velero Backup Configuration:

```
yaml
```

```
apiVersion: velero.io/v1
kind: Schedule
metadata:
  name: options-demo-backup
  namespace: velero
spec:
  schedule: "0 2 * * *" # Daily at 2 AM
  template:
    includedNamespaces:
      - options-demo
    includedResources:
      - persistentvolumeclaims
      - persistentvolumes
      - deployments
      - statefulsets
      - configmaps
      - secrets
    storageLocation: aws-s3
    volumeSnapshotLocations:
      - aws-ebs
    ttl: 720h # 30 days retention
```

## Disaster Recovery Runbook

**Recovery Time Objective (RTO):** 1 hour **Recovery Point Objective (RPO):** 24 hours

### Recovery Steps:

1. Restore OpenShift Cluster (via Terraform)
2. Apply Velero Restore
3. Verify Pod Health
4. Validate Data Integrity
5. DNS Cutover
6. Smoke Tests

---

## Security & RBAC

### OpenShift RBAC

#### Service Accounts:

```
yaml
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: options-demo-app
  namespace: options-demo
---
apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: options-demo-app-role
rules:
- apiGroups: [""]
  resources: ["configmaps", "secrets"]
  verbs: ["get", "list"]
- apiGroups: [""]
  resources: ["pods"]
  verbs: ["get", "list"]
```

## Network Policies:

```
yaml
```

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: app-network-policy
spec:
  podSelector:
    matchLabels:
      app: options-demo
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - podSelector:
            matchLabels:
              app: prometheus
      ports:
        - protocol: TCP
          port: 5000
  egress:
    - to:
        - podSelector:
            matchLabels:
              app: redpanda
      ports:
        - protocol: TCP
          port: 9092
```

## Secret Management

### Sealed Secrets für GitOps:

```
yaml

apiVersion: bitnami.com/v1alpha1
kind: SealedSecret
metadata:
  name: kafka-credentials
  namespace: options-demo
spec:
  encryptedData:
    username: AgBk7... (encrypted)
    password: AgCx9... (encrypted)
```

# Betriebskonzepte

## Update-Strategie

### Rolling Updates:

- Max Unavailable: 1
- Max Surge: 1
- Health Check Grace Period: 30s

### Canary Deployments:

1. Deploy 10% traffic to new version
2. Monitor for 15 minutes
3. Gradual rollout: 25% → 50% → 100%
4. Automatic rollback on error spike

## Performance Optimization

### Resource Tuning:

- Vertical Pod Autoscaler für Baseline
- Horizontal Pod Autoscaler für Traffic-Spikes
- Pod Disruption Budgets für Verfügbarkeit

### Cost Optimization:

- Node Autoscaling (2-8 Nodes)
- Spot Instances für Non-Critical Workloads
- Resource Requests basierend auf P95 Usage

## Incident Response

### Severity Levels:

- **SEV1:** Service Down, Customer Impact
- **SEV2:** Degraded Performance
- **SEV3:** Non-Critical Issues

### On-Call Playbook:

1. Alert Trigger
2. Initial Assessment (< 5min)
3. Incident Commander Assignment

4. Mitigation Steps
  5. Root Cause Analysis
  6. Post-Mortem
- 

## Lessons Learned & Best Practices

### DevOps-Mindset

#### Automatisierung:

- CI/CD Pipeline eliminiert manuelle Deployments
- Infrastructure as Code (Terraform) für Reproduzierbarkeit
- GitOps für Audit-Trail und Rollback-Fähigkeit

#### Dokumentation:

- Living Documentation (Code + Markdown)
- Runbooks für Operational Tasks
- Architecture Decision Records (ADRs)

#### Wissensteilung:




- Weekly Knowledge-Sharing Sessions
- Pair Programming für komplexe Tasks
- Open-Source Contributions


### Herausforderungen

1. **Kafka State Management:** StatefulSets mit Persistent Volumes
  2. **Monitoring Scale:** Prometheus Federation für Multi-Cluster
  3. **Log Volume:** Loki Retention Policies und Compaction
  4. **Security:** Pod Security Standards, Network Policies
- 

## Roadmap & Weitere Verbesserungen

### Phase 1 (Completed)





-  Docker Compose Entwicklungsumgebung
-  Options-Pricing Business-Logic
-  Monitoring mit Prometheus/Grafana

-  Logging mit Loki/Promtail

## Phase 2 (In Progress)

-  OpenShift Deployment
-  Helm Charts
-  Terraform IaC
-  GitOps mit ArgoCD

## Phase 3 (Planned)

-  Frontend Dashboard (React)
-  Service Mesh (Istio)
-  Multi-Cluster Federation
-  ML-basierte Alerting (AIOps)

---

## Anhang

### Nützliche Befehle

#### Lokale Entwicklung:

```
bash

# Stack starten
docker compose up -d

# Logs verfolgen
docker compose logs -f app

# Rebuild nach Code-Änderung
docker compose build app --no-cache
docker compose up -d
```

#### OpenShift Deployment:

```
bash
```

*# Cluster Info*

oc cluster-info

*# Deployment Status*

oc get pods -n options-demo

*# Logs anschauen*

oc logs -f deployment/options-demo-app

*# Port Forward für lokalen Zugriff*

oc port-forward svc/grafana 3000:3000

## Helm Operations:

bash

*# Install*

helm [install](#) options-demo ./helm -f values-prod.yaml

*# Upgrade*

helm upgrade options-demo ./helm -f values-prod.yaml

*# Rollback*

helm rollback options-demo

*# Debug*

helm template options-demo ./helm --debug

## Referenzen

- [OpenShift Documentation](#)
- [Helm Best Practices](#)
- [Prometheus Operator](#)
- [Loki Documentation](#)
- [GitOps with ArgoCD](#)

---

**Projekt Repository:** <https://github.com/AndrejLehner/OptionsDemo> **Autor:** Andrej Lehner **Version:** 1.0

**Datum:** Oktober 2025